

CS6910 Assignment 1

Lakshya J(EE19B035), Neham Jain(EE19B084), Nisharg Manvar(EE19B094)

March 2022

1 Function Approximation Task

In this task, we are asked to create a Multi Layer Feed Forward Neural Network of two layers where the Learning Rate and the number of neurons in each layer are hyper-parameters.

We are supposed to use Stochastic Gradient Descent as the optimizer and Mean Squared Error Loss for the loss function.

tanh was used as the activation function after every layer except for the last output layer.

I ran the model for 2000 epochs on batch sizes of 16, 32, 64, 128 for varying learning rates and number of neurons.

Learning rates used: [1e-6, 1e-5, 1e-4, 1e-3, 1e-2]

Number of neurons: [50, 100, 256, 512, 784, 1024, 2048]

Out of all these cases, I observed the least loss for the validation set in the case of batch size of 16, learning rate = 0.01 and number of neurons = 512.

Here is the result obtained:

Learning Rate : 0.01

No. of Neurons : 50

Validation loss : 0.0003136

I tested this model on the test set to obtain final results.

1.1 Plots Obtained

First, I plot the 3D plot of the actual function.

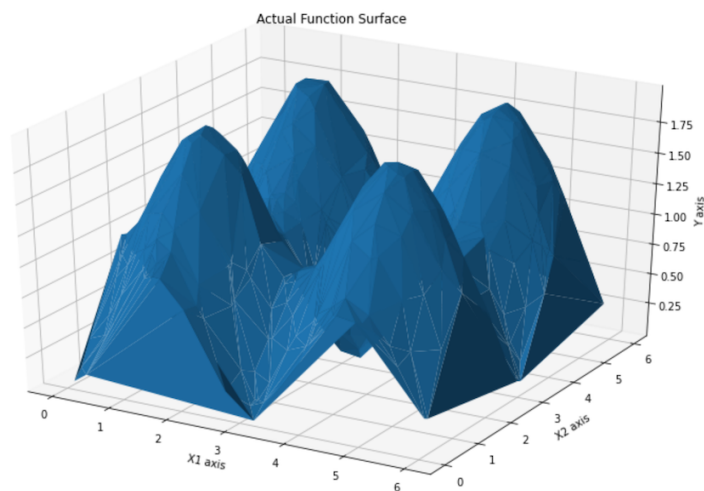


Figure 1: Actual Plot

Next I plot the output obtained after training the model for various epochs. I plot the output of the model for epochs 1, 2, 10, 50, 500, 1000, 1500 and 2000 (when the training is stopped).

It is observed that as the model moves through epochs, it is able to approximate the output better and give a smoother 3D output.




Figure 2: Plot of all data applied on the model after 1 epoch




Figure 3: Plot of all data applied on the model after 2 epoch

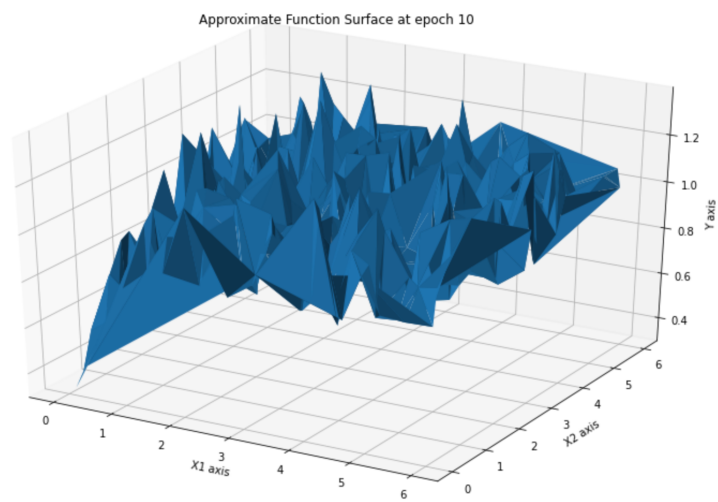


Figure 4: Plot of all data applied on the model after 10 epoch

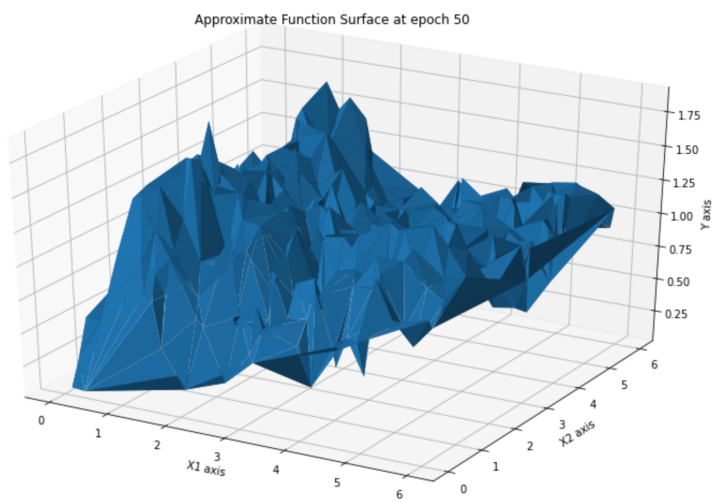


Figure 5: Plot of all data applied on the model after 50 epoch




Figure 6: Plot of all data applied on the model after 500 epoch




Figure 7: Plot of all data applied on the model after 1000 epoch




Figure 8: Plot of all data applied on the model after 1500 epoch




Figure 9: Plot of all data applied on the model after 2000 epoch

Next I plot the error plot for the t

1.2 Implementation Details:

To prevent the model from overfitting (blah blah blah)

training data, to see how well the training data learns with multiple epochs. We observe that the loss steeply drops within the first few epochs and then asymptotically reaches to a value after which it is not able to learn much, and oscillates around the found local optimum.



Figure 10: Training data error vs Number of epochs

The scatter plot after after training the model for 1500 epochs with the above parameters is as follows.

1.3 Observations

- As the number of epochs increases, the average error decreases. After a point, the output just oscillates around the local optimum.
- Train average loss: 0.00167
- Test average loss: 0.00114
- The scatter plot shows that the data points almost lie along the $x = y$ line, showing the model has learnt well.
- The final plot obtained at 2000 epochs almost resembles the actual function plot, except it is more sharper at the edges.




Figure 11: Scatter plot of actual output vs model obtained output of training data

- The validation loss varies with the epochs as follows:
Validation average loss at epoch 0000: 0.217057
Validation average loss at epoch 0100: 0.106351
Validation average loss at epoch 0200: 0.050664
Validation average loss at epoch 0300: 0.005119
Validation average loss at epoch 0400: 0.006670
Validation average loss at epoch 0500: 0.002086
Validation average loss at epoch 0600: 0.001214
Validation average loss at epoch 0700: 0.001714
Validation average loss at epoch 0800: 0.001274
Validation average loss at epoch 0900: 0.001025
Validation average loss at epoch 1000: 0.000955
Validation average loss at epoch 1100: 0.000678
Validation average loss at epoch 1200: 0.000656
Validation average loss at epoch 1300: 0.000648
Validation average loss at epoch 1400: 0.000698
Validation average loss at epoch 1500: 0.000476
Validation average loss at epoch 1600: 0.003656
Validation average loss at epoch 1700: 0.000401
Validation average loss at epoch 1800: 0.000398
Validation average loss at epoch 1900: 0.000314

- Final Train average loss: 0.00041
- Final Test average loss: 0.00031

2 Classifying Image Data

The images in the dataset given to us have been mapped to features of dimension 60. The features are then fed into our Multi Layer Feed Forward Neural Network model to classify the images into five distinct classes.

2.1 Hyper Parameter Optimization

We tune the

2.2 Implementation Details:

To prevent the model from overfitting (blah blah blah)

2.3 Results

2.4 Observations

training data, to see how well the training data learns with multiple epochs. We observe that the loss steeply drops within the first few epochs and then asymptotically reaches to a value after which it is not able to learn much, and oscillates around the found local optimum.